

Microsoft Warns Poisoned MCP Tool Descriptions Can Make AI Agents Leak Data

Swati Khandelwal Jun 30, 2026




New Microsoft research shows how attackers can hijack AI agents that act on a user's behalf, using nothing more than a poisoned tool description to make the agent quietly hand over company data to an outsider.

The Hacker News

June 2026

Researchers Build Self-Replicating AI Worm That Operates Entirely on Local, Open-Weight Models

Swati Khandelwal Jun 09, 2026



The Hacker News

University of Toronto researchers have built and tested a proof-of-concept AI-driven computer worm that uses a locally hosted open-weight large language model to reason its way through a network, generate tailored attack strategies for each target it encounters, and replicate itself, all without human intervention and without touching a commercial AI service.

June 2026